

Fine-Tuning Cross-Encoders for Detecting Hallucinations in Financial Customer Service Chatbots

Tahlee Stone

University of California, Berkeley, School of Information
tahlee.stone@berkeley.edu

Abstract

Mitigating hallucinations in Large Language Model (LLM) systems remains a key barrier to deployment in higher-risk domains such as finance and medicine where factual accuracy is both a regulatory requirement and component of responsible AI practice. This paper finetunes a cross-encoder DeBERTa-v3 transformer as a lightweight hallucination detection model to identify generated responses that are unsupported by or contradict retrieved evidence. The model is trained on Halubench benchmark data and a synthetic, domain-specific dataset simulating retail bank customer service chatbot responses. Its performance is evaluated against two baselines: a zero-shot prompted LLM-as-a-judge, based on the Lynx framework from Ravi et al. (2024), and a commercial AWS Bedrock contextual grounding guardrail. Results show sufficient performance for the cross-encoder model to be considered as a scalable, low latency component of hallucination detection in production-grade systems.

1 Introduction

Large Language Models (LLMs) are increasingly deployed in customer-facing systems, supporting tasks from dialogue generation to decision assistance (Brown et al., 2020; OpenAI, 2023). Despite their fluency and versatility, LLMs frequently produce hallucinations—outputs that are unverifiable or contradict retrieved evidence (Ji et al., 2023). In high-risk domains such as finance, these errors can erode user trust, and introduce compliance and reputational risks.

Retrieval-Augmented Generation (RAG) is a commonly used technique for mitigating hallucinations by incorporating externally retrieved documents into the model prompt to support response grounding (Lewis et al., 2020). However, hallucinations persist even in RAG settings, often due to incomplete retrieval, poor integration of evidence, or continued reliance on a model’s intrinsic knowl-

edge. As a result, post-generation hallucination detection has become an important component of production-grade LLM systems, functioning as a safeguard or ‘guardrail’ before output delivery.

This paper aims to fine-tune a lightweight cross-encoder on benchmark and synthetic financial data for hallucination detection. It compares this model against an open-source LLM-as-a-judge and a commercial grounding guardrail to assess its viability as a low-latency, scalable guardrail in production LLM-based systems for higher-risk domains like finance.

Hallucination detection is typically framed as a *groundedness classification* task: for a given query, retrieved context, and model-generated response triplet, the guardrail must determine whether the response is supported by the evidence provided (Luo et al., 2024). Effective solutions must handle semantic nuance, domain-specific language, and variable response structures.

1.1 Background

Several groundedness evaluation frameworks have emerged to address this challenge, with LLM-based approaches among the most widely adopted (Deng et al., 2023; Luo et al., 2024). **RAGAS** (Es et al., 2023), an open-source toolkit based on heuristic metrics, evaluates answer faithfulness by prompting an LLM to extract factual claims from a response and verify each claim against the retrieved context, returning a score based on the proportion of grounded claims.

LLM-as-a-judge methods, including both prompted and fine-tuned models, assess hallucinations by reasoning over the full query, context, and response to provide a more holistic assessment of groundedness. Prompted approaches use general-purpose models such as GPT-4 to classify responses relative to retrieved context (Zheng et al., 2023; Manakul et al., 2023), while fine-tuned LLM-as-a-judges such as the open-source **Lynx** (Ravi

et al., 2024) model are explicitly trained on hallucination datasets. Lynx has demonstrated stronger performance than GPT-4 on QA-style hallucination detection tasks.

Some commercial guardrail systems—e.g., [AWS Bedrock’s contextual grounding](#) guardrail—offer API-based hallucination checks for LLM outputs, but their inner mechanisms are opaque and they may introduce latency, cost, and privacy risks.

While performant, LLM-based evaluators present several practical limitations: they are costly to run at scale, introduce latency, and produce non-deterministic outputs—where repeated evaluations of the same input can yield different results. This limits reproducibility, complicates auditing, and creates inconsistencies in quality assurance workflows. These challenges motivate the need for efficient, locally deployable alternatives.

Encoder-based models offer a competitive alternative to LLM-based hallucination detection, combining transparency, low inference overhead, and real-time feasibility. **Luna** is a DeBERTa-large cross-encoder fine-tuned on multi-domain RAG data to detect hallucinated spans at the token level (Belyi et al., 2024). It outperforms GPT-3.5-based evaluators on the RAGTruth benchmark while achieving up to 97% cost reduction and 91% lower latency. **Lettuce Detect**, built on the ModernBERT architecture, surpasses Luna and GPT-4-turbo with a 79.2% F1—despite being 30× smaller than top LLMs. (Ádám Kovács and Recski, 2025).

1.2 Contributions

This paper builds on recent work by addressing key limitations of LLM-based hallucination detection, and proposes a lightweight cross-encoder classifier based on DeBERTa-v3 for groundedness classification. The model is fine-tuned on a combination of benchmark data and a synthetic financial-domain dataset simulating retail bank Q&A customer service interactions. Its performance is evaluated against two LLM-based baselines: an open-source zero-shot verifier based on the Lynx framework Ravi et al. (2024), and a closed-source commercial grounding guardrail provided by [AWS Bedrock](#). The contributions of this paper are:

1. Develop a practical, resource-efficient cross-encoder approach for hallucination detection in RAG systems that balances performance and latency;

2. Construct a domain-specific evaluation set reflecting the language and groundedness challenges of retail bank dialogue systems, leveraging the [Banking77 dataset](#) to enrich task realism and intent coverage¹;
3. Conduct a comparative evaluation of the cross-encoder against both open-source and commercial LLM-based evaluators, highlighting trade-offs in groundedness accuracy, transparency, and deployability.

Together, these findings offer actionable guidance for operationalising hallucination detection in higher-risk environments such as banking and finance, where LLM outputs must be both reliable and auditable.

2 Methodology

2.1 Data Construction

The training and evaluation data were sourced from two key datasets: (i) the HaluBench benchmark dataset; and (ii) a synthetic dataset of customer service queries for a retail-banking chatbot.

HaluBench Benchmark

This benchmark dataset contains around 15,000 query–context–response triplets annotated for hallucination detection. Introduced by Ravi et al. (2024) to develop the LYNX model, HaluBench combines multiple established datasets from different domains.² It includes both naturally occurring hallucinations generated by large language models and synthetic perturbations. The benchmark consists of the following domain-specific subsets:

- **FinanceBench**—financial reasoning (Islam et al., 2023)
- **DROP**—numerical reasoning (Dua et al., 2019)
- **PubMedQA**—biomedical question answering (Jin et al., 2019)
- **RAGTruth**—retrieval-augmented generation evaluation (Wu et al., 2023)

Two sub-datasets—COVID-QA (Möller et al., 2020) and HaluEval (Li et al., 2023) datasets—were

¹<https://huggingface.co/datasets/PolyAI/banking77>

²See Section 3.2 of Ravi et al. (2024) for further detail on the construction of the HaluBench benchmark dataset.

excluded from analysis due to lengthy report-style contexts and limited domain relevance, and due to a systematic length bias between grounded and hallucinated responses, respectively.

A final set of 3,900 observations were sampled from Halubench for both fine-tuning and evaluation (with 3,315 for training and 585 for the hold-out evaluation set).³

RetailbankQA Synthetic Dataset

To fine-tune a hallucination detection model in the context of a retail bank customer service setting, a synthetic RetailbankQA dataset was constructed using a multi-stage prompting pipeline shown in Figure 1.

- **Step 1: Query Sampling.** A representative sample of 1,385 customer queries were obtained from the [Banking77 dataset](#), covering 77 intent categories (e.g., card_payment_problems, atm_support, interest_rate_query).⁴
- **Step 2: Context Generation.** For each query from Step 1, gpt-4o-mini was prompted to generate a 150–250 word FAQ-style passage that could plausibly be used to answer the customer’s question in a banking customer service setting.
- **Step 3: Response Generation.** Two responses were created for each query-context pair:
 - **Grounded responses** were generated using gpt-4o with temperature set to zero, and prompted to rely strictly on the accompanying context passage, without introducing any external knowledge.
 - **Hallucinated responses** were generated using gpt-3.5-turbo-0125, with prompting designed to subtly introduce factual inconsistencies (e.g., fabricated service features or incorrect fees) while maintaining surface fluency and plausible relevance.

³The reduction in sample size compared to the full HALUBENCH dataset reflects the exclusion of the HALUEVAL subset (approximately 10,000 examples), which was removed due to a systematic length bias between hallucinated and grounded responses.

⁴The Banking77 dataset contains 13,083 real-world customer service queries from the banking domain. It is widely used for research in intent classification and banking-domain conversational AI.

- **Step 4: Dataset Assembly.** The final dataset included 2,770 labeled (query, context, response) triplets—1,385 grounded and 1,385 hallucinated. Of these, 1,000 grounded–hallucinated pairs were allocated for training, and 385 pairs were held out as an evaluation set. The train–validation split was performed using unique pair identifiers to ensure no query–context–response overlap across splits, thereby preventing data leakage. All examples were formatted consistently with the HaluBench schema to support benchmark-compatible training and evaluation workflows.

- **Step 5: LLM-based Annotation Verification.** To assess label quality, a random sample of 50 examples (25 grounded, 25 hallucinated) from the synthetic dataset was reviewed by OpenAI’s o3-pro model⁵, prompted to act as an expert human QA reviewer and annotator. This model offers stronger reasoning capabilities than those used during data construction (gpt-3.5-turbo-0125 for hallucinations; gpt-4o-mini for context generation; and gpt-4o for grounded responses). Following strict groundedness criteria, the annotation process showed a 90% agreement with the original labels. This degree of alignment suggests that the synthetic labels are broadly consistent with expert judgment and sufficiently reliable for training purposes.⁶

Final sample

The final sample consists of 6,600 observations, with 3,900 from HaluBench and 2,770 test cases from the constructed RetailbankQA dataset.

Domain	# Examples	%H	Med CT	Max CT
RetailBankQA	2,770	50%	275	352
HaluBench	3,900	34%	291	2494
DROP	1,000	50%	224	1204
FinanceBench	1,000	50%	454	2494
PubMedQA	1,000	50%	280	937
RAGTruth	900	18%	282.5	538

Table 1: Summary statistics of full training and evaluation dataset, where %H denotes share of hallucinated responses and CT indicates context token length. HaluBench includes benchmark datasets from finance, biomedical, and general knowledge domains.

⁵<https://platform.openai.com/docs/models/o3>

⁶Disagreements primarily involved borderline cases with minor omissions or over-specificity, underscoring the need for clearer prompting or light label smoothing in future iterations.




Figure 1: Synthetic data construction pipeline showing the four stages: sampling customer queries, generating synthetic contexts, creating grounded and hallucinated responses, and producing the final labeled dataset.

Table 1 reports token-level statistics across all datasets. The synthetic retail bank dataset (Retail-bankQA) contains moderate context lengths compared to other datasets in HaluBench. The share of hallucination are balanced in all datasets but RAGTruth. Each dataset was split using stratified sampling into training, validation, and test subsets (70/15/15) to preserve class balance. Validation sets were used for hyperparameter tuning, and final evaluation was conducted on held-out test sets. All examples were tokenized using the DeBERTa-v3-base tokenizer.

3 Model

The model used for binary hallucination classification is a fine-tuned cross-encoder based on the **DeBERTa-v3-base** architecture.⁷ This architecture has demonstrated strong performance on natural language inference (NLI) benchmarks such as MNLI, ANLI, and SNLI, outperforming models like BERT, RoBERTa, and XLNet (He et al., 2021; Liu et al., 2019; Yang et al., 2019).

DeBERTa-v3-base is well-suited for hallucination detection tasks that require fine-grained factual alignment, as it is pretrained on entailment-style objectives. Compared to bi-encoders or prompt-based LLM evaluators, the cross-encoder architecture jointly encodes the *query*, *retrieved context*, and *response*, enabling the model to capture rich inter-segment dependencies. This design is consistent with recent hallucination detection frameworks such as Luna (Belyi et al., 2024) and Lettuce Detect

⁷<https://huggingface.co/microsoft/deberta-v3-base>

(Ádám Kovács and Recski, 2025), and aligns with answer faithfulness metrics used by LLM evaluation protocols.

3.1 Training Setup

To adapt the pretrained model for binary hallucination detection, a classification head is added on top of the DeBERTa encoder to produce a single logit representing the likelihood that the response is grounded. The input is formatted as:

[CLS] query [SEP] context [SEP] response [SEP]

The model is trained using the Hugging Face Trainer API with a batch size of 16, a learning rate of 2×10^{-5} , and AdamW optimizer. A linear learning rate scheduler with warm-up is used. The loss function is **binary cross-entropy**, with labels mapped as 1 for grounded (PASS) and 0 for hallucinated (FAIL). Training is conducted over three epochs on a balanced dataset comprising examples from both HaluBench (general domain) and Retail-BankQA (financial domain). The Hugging Face DeBERTa tokenizer is used to encode the input sequence.

3.2 Long Context Handling

To handle long inputs exceeding the 512-token limit, a **windowed chunking strategy** is applied. Each input is split into overlapping chunks and passed individually through the model. Inference scores across chunks are then aggregated to produce an **example-level groundedness prediction**. This approach draws inspiration from long-context RAG evaluation strategies introduced in Luna (?),

Dataset	Model	Accuracy	Macro Precision	Macro Recall	Macro F1
HaluBench	DeBERTa-v3 Cross-encoder	69.0%	70.0%	68.0%	68.0%
	GPT-4o (LLM Evaluator)	82.6%	82.9%	83.6%	82.5%
	AWS Bedrock	58.6%	60.5%	60.3%	58.6%
- DROP	DeBERTa-v3 Cross-encoder	88.0%	85.9%	89.7%	87.8%
	GPT-4o (LLM Evaluator)	82.4%	82.7%	82.2%	82.2%
	AWS Bedrock	47.9%	46.7%	47.1%	45.8%
- FinanceBench	DeBERTa-v3 Cross-encoder	52.2%	51.9%	75.6%	61.5%
	GPT-4o (LLM Evaluator)	84.8%	86.1%	84.9%	84.7%
	AWS Bedrock	50.6%	51.1%	50.8%	47.3%
- RAGTruth	DeBERTa-v3 Cross-encoder	75.8%	85.1%	85.1%	85.1%
	GPT-4o (LLM Evaluator)	84.8%	74.8%	77.2%	75.9%
	ASW Bedrock	80.8%	69.9%	74.8%	71.7%
- pubmedQA	DeBERTa-v3 Cross-encoder	68.6%	67.0%	82.9%	74.1%
	GPT-4o (LLM Evaluator)	77.9%	82.0%	79.4%	77.6%
	AWS Bedrock	60.0%	62.2%	61.3%	59.6%
RetailbankQA	DeBERTa-v3 Cross-encoder	97.0%	97.0%	97.0%	97.0%
	GPT-4o (LLM Evaluator)	79.0%	85.2%	79.0%	78.0%
	AWS Bedrock	80.9%	85.5%	80.9%	80.3%

Table 2: Evaluation results across six datasets comparing cross-encoder, LLM-based, and AWS Bedrock hallucination detection methods. Results are reported as Accuracy, Macro Precision, Macro Recall, and Macro F1. RetailbankQA shows the largest gains from the cross-encoder approach.

which emphasize the importance of distributed support across large retrieved contexts.

While Luna performs hallucination detection at the span level using token-level support probabilities, this implementation focuses on **example-level binary classification**. This design choice trades off granularity for efficiency, enabling faster inference and simpler integration into applications such as chatbot QA systems, offer ranking engines, and safety filters in high-risk domains like finance.

To align with the 512-token input limit of DeBERTa-v3-base, token lengths were analyzed (see Figures 2 and ??). Most examples fell below this threshold. For the small fraction exceeding the limit (<5%), context sequences were truncated to 400 tokens, with the full response retained to preserve grounding cues.




Figure 2: Distribution of context token length across all datasets (dashed line = 512-token max input for DeBERTa-v3-base).




Figure 3: Performance comparison of cross-encoder versus LLM-based models; by evaluation metric.

4 Baseline Comparison Models

Baseline comparisons include a prompted GPT-4o evaluator and a commercial grounding guardrail from AWS Bedrock.

4.1 Prompted LLM Verifier

A zero-shot evaluation baseline is implemented using OpenAI’s gpt-4o-2024-05-13 model, accessed via the OpenAI API.⁸ The prompting approach is adapted from the Lynx framework proposed by Ravi et al. (2024) and is presented using a structured Jinja2 template. The full prompt is included in Appendix A. Each input is formatted as a (query, document, response) triplet, and the

⁸<https://platform.openai.com/docs/models/gpt-4o>

model is instructed to assess whether the response is grounded in the retrieved document, based on strict faithfulness criteria.

The model returns a JSON object containing a REASONING field (a brief justification) and a SCORE field, which is parsed as either "PASS" or "FAIL". These binary outputs are used for groundedness classification in the evaluation.

4.2 AWS Bedrock Guardrail

As an off-the-shelf comparator model, the commercial contextual grounding guardrail from Amazon Bedrock was selected as a closed-source baseline.⁹ Each (query, context, response) triplet is submitted via the apply_guardrail API with contextual grounding policies enabled and relevance checks disabled. The system uses Claude Sonnet 4, the foundation model available via AWS at the time of evaluation, to return a grounding score and a threshold-based decision.

A fixed threshold of 0.7 is applied to determine whether the response is grounded. Outputs are mapped to binary PASS or FAIL labels to support direct comparison with other evaluators. Although the guardrail is configured to block ungrounded responses, its decision is used solely for classification in this study. This baseline was selected for its alignment with the contextual grounding task and its accessibility as a managed service.

4.3 Evaluation

All models are evaluated on both datasets using gold-labeled ground truth. We report accuracy, precision, recall, and F1 scores for both PASS and FAIL classes, as well as macro-averaged metrics. Confusion matrices and qualitative failure cases are also analyzed to understand model behavior under domain shift and subtle hallucination patterns.

5 Results

We evaluate groundedness detection performance across two datasets: the HaluEval benchmark and a domain-specific synthetic challenger set. Table ?? reports accuracy, class-wise F1 scores, and macro-averaged F1 for each model.

5.1 Key Findings

- On the HaluEval benchmark, the LLM verifier slightly outperforms Bedrock across all

⁹<https://docs.aws.amazon.com/bedrock/latest/userguide/guardrails-contextual-grounding-check.html>

QUESTION
How to explain quartiles?
REFERENCE
Passage 2: Second quartile (Q2), which is more commonly known as the median, splits the data in half (50%). The median divides the data into a lower half and an upper half. Third quartile (Q3), also known as the upper quartile, splits the lowest 75% (or highest 25%) of data.
RESPONSE
The first quartile (Q1) splits the lowest 25% of the data, while the second quartile (Q2) splits the data into two equal halves, with the median being the middle value of the lower half. Finally, the third quartile (Q3) splits the highest 75% of the data.
ANNOTATION
Span: highest 75%
Type: Evident Conflict
Reason:
Original: "lowest 75% or highest 25%"
Generative: "highest 75%"

Table 3: Example of a RAGTruth data point, including the question, reference, model-generated response, and annotated conflict span.

metrics, particularly in precision and macro F1.

- On the domain-specific challenger dataset, Bedrock outperforms the LLM verifier in overall accuracy and macro F1, indicating stronger generalization to narrow financial contexts.
- F1 scores for hallucinated (FAIL) responses drop significantly in the challenger set for both models. The LLM verifier shows a notable drop to 0.734, suggesting reduced sensitivity to subtle domain-specific inconsistencies.

6 Discussion

The results highlight several key trade-offs between system types:

LLM-based verification. GPT-4o achieves the highest overall performance on the HaluEval benchmark, particularly in identifying grounded responses (Recall_PASS = 0.934). However, it is more sensitive to subtle hallucinations in the financial domain, with reduced performance on the challenger set (F1_FAIL = 0.734).

AWS Bedrock. The groundedness guardrail from AWS Bedrock is surprisingly competitive, with strong F1_FAIL on both datasets (0.774 on HaluEval, 0.767 on Challenger). However, it shows a higher rate of false positives on grounded responses, indicating a tendency toward over-conservatism in its classification decisions.

Cross-encoder model. The DeBERTa-v3 cross-encoder demonstrates a strong ability to balance fine-grained semantic detection with efficiency. Its classification head, fine-tuned on domain-specific data, effectively captures grounding inconsistencies that general-purpose models often miss. This enables both high accuracy and better generalization to financial dialogue.

Overall takeaway. These findings reinforce the limitations of relying solely on zero-shot LLM or commercial guardrails in domain-specific RAG systems. Fine-tuned cross-encoders offer a practical middle ground—delivering robust groundedness classification with lower latency and improved control, making them a promising component for scalable LLM safety pipelines.

`\usepackage{graphicx}.`

7 Results and discussion

7.1 Main Evaluation

7.2 Error Analysis

[Discuss types of hallucinations missed, false positives, patterns across datasets.]

8 Conclusion

[Summary of findings. Implications for scalable, reliable financial dialogue systems. Future work.]

Limitations

References

- Masha Belyi, Robert Friel, Shuai Shao, and Atindriyo Sanyal. 2024. [Luna: An evaluation foundation model to catch language model hallucinations with high accuracy and low cost](#).
- Tom B Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, et al. 2020. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901.
- Zhihong Deng et al. 2023. [A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions](#). *arXiv preprint arXiv:2311.05232*.
- D. Dua, Y. Wang, P. Dasigi, G. Stanovsky, S. Singh, and M. Gardner. 2019. Drop: A reading comprehension benchmark requiring discrete reasoning over paragraphs. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 2368–2378.
- Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. 2023. [Ragas: Automated evaluation of retrieval augmented generation](#). *arXiv preprint arXiv:2309.15217*.
- Pengcheng He, Xiaodong Gao, Jianfeng Chen, and Jianfeng Gao. 2021. Deberta: Decoding-enhanced bert with disentangled attention. *arXiv preprint arXiv:2006.03654*.
- M. Islam, Y. Zhang, and D. Lee. 2023. Financebench: A financial reasoning benchmark for llms. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1012–1023.
- Zhewei Ji, Nayeon Lee, Jason Fries, et al. 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys*.
- Wenhao Jin, Yifan Peng, Luo Si, and Xuan Li. 2019. Pubmedqa: A dataset for biomedical research question answering. In *Proceedings of the 2019 Annual Meeting of the Special Interest Group on Bioinformatics of the ACL (BioNLP-ACL)*, pages 256–265.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*.
- Junyi Li, Xiaoxue Cheng, Wayne Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. 2023. Halueval: A large-scale hallucination evaluation benchmark for large language models. *arXiv preprint arXiv:2305.14269*.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*.
- Junliang Luo, Tianyu Li, Di Wu, Michael Jenkin, Steve Liu, and Gregory Dudek. 2024. [Hallucination detection and hallucination mitigation: An investigation](#). *arXiv preprint arXiv:2401.08358*.
- Potsawee Manakul, Adian Liusie, and Mark JF Gales. 2023. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. *arXiv preprint arXiv:2303.08896*.
- Philip Möller, Alexander Yates, Qian Li, and Jiawei Xu. 2020. Covid-qa: An open domain question answering dataset for covid-19. In *Proceedings of the 2020 ACL Workshop on NLP for COVID-19*, pages 30–40.
- OpenAI. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Selvan Sunitha Ravi, Bartosz Mielczarek, Anand Kannappan, Douwe Kiela, and Rebecca Qian. 2024. [Lynx: An open source hallucination evaluation model](#). *arXiv preprint arXiv:2407.08488*.

Yi Wu, Liang Chen, and Hector Fernandez. 2023. Ragtruth: A benchmark for evaluating retrieval-augmented generation. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, pages 8450–8462.

Zhilin Yang, Zihang Dai, Yiming Yang, Jaime Carbonell, Ruslan Salakhutdinov, and Quoc V Le. 2019. Xlnet: Generalized autoregressive pretraining for language understanding. *Advances in neural information processing systems*, 32.

Yuxuan Zheng, Weilu Zheng, Diyi Wang, Bill Yuchen Lin, and Xiang Ren. 2023. Judging groundedness in llms: A judge model approach. In *Findings of the Association for Computational Linguistics: EMNLP 2023*.

Ádám Kovács and Gábor Recski. 2025. [Lettucedetect: A hallucination detection framework for rag applications](#). *arXiv preprint arXiv:2502.17125*. Version 1, submitted 24 Feb 2025.

```
{"REASONING": <brief reasoning as bullet points>,  
"SCORE": <your final score>}
```

A Groundedness Evaluation Prompt

The following prompt was used for LLM-based groundedness evaluation (GPT-4o). Query, context, and response fields were dynamically filled for each example.

Given the following QUESTION, DOCUMENT and ANSWER you must analyze the provided answer and determine whether it is faithful to the contents of the DOCUMENT.

A faithful ANSWER must not offer new information beyond the context provided in the DOCUMENT.

A faithful ANSWER must be grounded in the DOCUMENT and not based on external knowledge.

A faithful ANSWER also must not contradict information provided in the DOCUMENT.

Even if there is a slight discrepancy between the ANSWER and the DOCUMENT, or any ungrounded assumptions made in the ANSWER, the ANSWER must be considered not faithful.

The faithfulness of the ANSWER is independent of its relevancy to QUESTION.

Even if the ANSWER is irrelevant to the QUESTION, your task is to evaluate if it is faithful to the DOCUMENT regardless of relevancy.

Output your final verdict by strictly following this format: "PASS" if the answer is faithful to the DOCUMENT and "FAIL" if the answer is not faithful to the DOCUMENT. Briefly, show your reasoning.

—
QUESTION (THIS DOES NOT COUNT AS BACKGROUND INFORMATION):

{{user_query}}

—
DOCUMENT:

{{retrieved_context}}

—
ANSWER:

{{model_answer}}

—
Your output should be in **JSON format** with the keys "REASONING" and "SCORE":